

Advanced Methods in Text Analytics

Transformers



Why Transformers?

- **Most successful and commonly used** deep learning architecture in NLP
 - By far!
- **Resilient technology**
 - RNNs improved over FNNs
 - LSTMs improved over vanilla RNNs
 - Attention improved over vanilla encoder-decoder architectures
 - Transformer proposed in 2017/2018, still largely unchanged
 - Alternative architectures do exist/are proposed, e.g. [Mamba](#), [LSTMs](#)
 - But transformers difficult to replace, lots of infrastructure in place
 - Improvements must be worth the effort
- Main advantages over RNNs (previous state-of-the-art models)
 - **Attention over arbitrarily large inputs**, i.e. no recurrent connections
 - **Flexible/powerful attention** mechanism
 - **Scalable!** Parallelizable, important contrast to RNN-based models

Outline


1. Recap: Attention

2. Self-Attention

3. The Transformer Architecture

Recap: Attention

Encoder-Decoder Architecture




- **Goal:** create contextually appropriate sequence of arbitrary length
 - Known as seq2seq models
- Components
 - **Encoder:** typically, an RNN
 - **Context vector:** produced by encoder (typically last hidden state in RNN)
 - **Decoder:** RNN, produces task-dependent output based on context vector
 - Thus, input sequence represented entirely by context vector
- **Common applications:** machine translation, dialogue systems, etc.

Attention (1)

- Seminal work in machine translation: [Bahdanau et al. \(2015\)](#)
- Recall: RNNs have a hard time using information far back in time
 - Thus, the **context vector may not encode everything we need**
 - LSTMs somewhat address this, but...
- **Why not allow decoder to access input sequence?**
- In encoder-decoder architecture:
 - Decoder accesses input via hidden states of encoder
 - **Context vector** is now **weighted sum of hidden states in encoder**
- **Important:** context vector dependent on decoder state!
 - Thus, model “attends to” different parts of input to produce different parts of output
- Let’s look at all of this in more detail!


Attention (2)

- $\mathbf{c}_i$ is context vector for output i , defined as $\mathbf{c}_i = \sum_{j=1}^n \alpha_{ij} \mathbf{h}_j^e$
 - That is, $\mathbf{c}_i$ is a weighted sum of encoder hidden states
- Coefficients α_{ij} known as **attention weights**
 - Encode **how much attention is paid to input j** to produce output i
 - Decoder output $\mathbf{h}_{i-1}^d$ necessary, encodes output $i-1$ (dashed lines)



Attention (3)

- Entry i, j corresponds to attention weight for input i given target token j
- We can see most words are translated 1-to-1, i.e. look at n th input word to produce n th output word
- But some are not so simple!
- To produce word *européenne*, the 7th output word, model looks at 5th input word, the actual relevant one!
- **Relevant input can be far!**
 - “Ich habe heute Abend Bratwurst mit Bröt und Kartoffeln *gegessen*.”
 - “Tonight I *ate*...”




How to Compute Attention Weights?

- Two steps:
 1. Compute **attention scores**
 2. Combine attention scores to produce **attention weights**
- **Attention scores:** how relevant each input is to decoder's last state
 - Each input $\rightarrow$ hidden states in encoder $\mathbf{h}_j^e$
 - Each output $\rightarrow$ decoder hidden states $\mathbf{h}_i^d$
- Common approach: **dot-product attention**
 - $score(output\ i, input\ j) = (\mathbf{h}_i^d)^T \mathbf{h}_j^e$
 - Used to compute output $i + 1$
 - Relevance seen as similarity
- Similarities computed between each input token and last decoder state
 - We still don't know the *relative* relevance across input words
- **Attention weights:** encode *relative* relevance across input words
 - $\alpha_{ij} = softmax(score(output\ i, input\ j))$, i.e. mass of α_{ij} w.r.t. all other $\alpha_{ij'}$

Attention: In Short

- To produce output $i+1$:
 1. Compute **attention scores** using last hidden state of decoder's $\mathbf{h}_i^d$, e.g. with dot-product attention: $score(output\ i, input\ j) = (\mathbf{h}_i^d)^T \mathbf{h}_j^e$
 2. Compute **attention weights**, e.g. using softmax across all attention scores
 3. Compute **context vector** $\mathbf{c}_i$ as linear combination of encoder hidden states, where coefficients are attention weights, i.e. $\mathbf{c}_i = \sum_j \alpha_{ij} \mathbf{h}_j^e$
 4. Use $\mathbf{c}_i$, along with y_i and $\mathbf{h}_i^d$, to produce output y_{i+1}



- Thus, attention allows a seq2seq model to see *dynamic* representation of input at each output step
 - Different operations can be used to compute scores and weights

Self-Attention

The Heart of Transformers

- Transformers were introduced by [Vaswani et al.](#) in 2017
- Virtually all of state-of-the-art NLP is based on transformers, e.g.
 - **BERT** and other **masked language models** that provide word representations (covered soon)
 - The **GPT family** of text-generating language models (covered soon)
 - State-of-the-art text encoders used for retrieval (see our IR course)
- Deep architecture with many components
- **Key component** of the transformer architecture: **self-attention**
 - They proposed to drop the RNNs and “simply” focus on attention
 - In reality, more components aside from attention
 - But self-attention is **main innovation**, allowed for **flexible and scalable attention over long sequences**
 - More on this later
- First, let’s focus on self-attention
 - Then on the overall architecture of the model

Attention as Information Retrieval


- **Useful intuition** about attention **comes from information retrieval**
- Say you open YouTube and input the **query** “*cats dressed as Batman*”
- The search system represents each video with a set of **keys**
 - Think attributes, class properties
 - For example: *title, description, channel_name, publication_date*, etc.
- For any query-key match the system finds, it returns a **value**
 - In this case, values are relevant videos, e.g. those with text similar to your query in the title, description, etc.
- We can see attention as such a query-keys-values system
 - Given attention score $score(\text{output } i, \text{input } j) = (\mathbf{h}_i^d)^T \mathbf{h}_j^e$, decoder state $\mathbf{h}_i^d$ is a **query**, encoder state $\mathbf{h}_j^e$ is a **key**
 - Context vector $\mathbf{c}_i$ represents retrieved values, but as linear combination?
 - Yes, so $\mathbf{c}_i = \sum_j \alpha_{ij} \mathbf{h}_j^e$ where $\mathbf{h}_j^e$ are again used as values
- Thus, attention can be seen as a *soft-retrieval* system
 - It returns relevant values, but each weighted by how relevant they are

Self-Attention

- **Attention:** compare item of interest to collection of *other items* in a way that reveals their relevance in the current context
 - In encoder-decoder architecture, item of interest is *decoder* state, i.e. output sequence so far
 - Collection of other items is hidden states of *encoder*, i.e. input sequence
- **Self-attention:** compare each token in given sequence, to all other tokens in the *same sequence*, i.e. item of interest is in same collection of items to compare with
- Thus, we use **representations from the same sequence as queries, keys and values**
 - $score(input\ i, input\ j) = (\mathbf{x}_i^e)^T \mathbf{x}_j^e$
 - $\alpha_{ij} = softmax(score(input\ i, input\ j))$
 - $\mathbf{y}_i = \sum_j \alpha_{ij} \mathbf{x}_j^e$where $\mathbf{x}_j$ are representations of given sequence, and $\mathbf{y}_i$ are representations after attention, i.e. **contextualized** representations


Visualizing Self-Attention

- **Intuition of self-attention:** given word, we *pay attention* to its context (i.e. surrounding words) to determine its meaning
 - Some words in context may be better signal for meaning than others
 - E.g. word **reading** likely **more important than** word **always** to determine meaning of word **book** in sentence "*I always enjoy reading a book.*"
- **Visualization of self-attention layer:** given sequence of word vectors, layer produces same sequence but of **contextualized word vectors**



Self-Attention Layer (1)

- **Input:** sequence of n tokens
- **Output:** sequence of n contextualized tokens




- Layer is **parameterized by** matrices W^Q , W^K and W^V
 - Each a linear transformation applied to input tokens x_i
 - **Each transformation** is applied when tokens are used in the different roles: **queries**, **keys** and **values**

Self-Attention Layer (2)

- Note the contrast to attention so far, which was not parameterized!
 - Attention scores: dot-product
 - Attention weights: softmax
- Before, model had to learn token representations that could be combined to provide useful context vectors
 - The **parameters in self-attention** allow for a **more flexible attention mechanism**
- Specifically, $\mathbf{q}_i = \mathbf{W}^Q \mathbf{x}_i$, $\mathbf{k}_i = \mathbf{W}^K \mathbf{x}_i$ and $\mathbf{v}_i = \mathbf{W}^V \mathbf{x}_i$ for input $\mathbf{x} \in R^{D \times 1}$
- Then:
 - $score(input\ i, input\ j) = \mathbf{q}_i^T \mathbf{k}_j$
 - $\alpha_{ij} = softmax(score(input\ i, input\ j))$
 - $\mathbf{y}_i = \sum_j \alpha_{ij} \mathbf{v}_j$
- Note that for dot product attention, we require that $\mathbf{W}^Q, \mathbf{W}^K \in R^{D' \times D}$
 - In general, $\mathbf{W}^Q$, $\mathbf{W}^K$ and $\mathbf{W}^V$ can be of any size allowed by the required computations for scores and weights

Self-Attention Layer (3)

- Here, model attends only to previous tokens
 - **Design choice, depends on task**
- Note:
 - **Transformations** to k , q and v
 - **Dot products** with one q and each k to compute scores
 - **Softmax** to compute weights
 - **Linear combination** with each v to get contextualized representation y




Self-Attention Layer (4)

- The original model actually used scaled dot-product attention
 - $score(input\ i, input\ j) = \mathbf{q}_i^T \mathbf{k}_j / \sqrt{D_k}$ where D_k is size of keys and query vectors
 - Done to avoid numerical stability issues when using softmax
- We can represent the computation for all $\mathbf{y}_i$ as matrix products
 - Let $\mathbf{X} \in \mathbb{R}^{N \times D}$ be the matrix of N input tokens of size D
 - Then: $\mathbf{Q} = \mathbf{XW}^Q \in \mathbb{R}^{N \times D'}$, $\mathbf{K} = \mathbf{XW}^K \in \mathbb{R}^{N \times D'}$, $\mathbf{V} = \mathbf{XW}^V \in \mathbb{R}^{N \times D}$
 - All relevant dot products, i.e. attention scores, are in $\mathbf{QK}^T \in \mathbb{R}^{N \times N}$
 - Apply row-wise softmax for attention weights: $softmax(\mathbf{QK}^T / \sqrt{D_k}) \in \mathbb{R}^{N \times N}$
 - Then multiply by $\mathbf{V}$ to get final stacked contextualized representations
- All together: $Self-Attention(\mathbf{Q}, \mathbf{K}, \mathbf{V}) = softmax(\mathbf{QK}^T / \sqrt{D_k}) \mathbf{V}$
- **Cost:** each $\mathbf{y}_i$ computed independently, thus highly parallelizable
 - But attention quadratic in input size (every item with every other item)
 - Thus, size of input sequence a fundamental limitation of this architecture
 - **Runtime increases quadratically with size of input sequence**

The Transformer Architecture

The Transformer Architecture

- Originally an encoder-decoder architecture
- **Encoder:** multi-head **attention**
 - Plus additional components
- **Decoder:** multi-head **attention**
 - Plus additional components
- **Interaction between encoder/decoder:**
 - Multi-head **attention**
- Hence the name: **Attention is All you Need**
- **Multi-head attention:** multiple self-attention layers (discussed soon)
- Additionally, **other components:**
 - Positional encoding
 - Layer normalization
 - Residual connections
 - FNN
- Let's look at them in more detail, but first...




The Animated Transformer

- From [Google's blog](#), here's how representations are built/interact
 - 3 encoder layers are stacked for depth (still deep learning)
 - 3 decoder layers are also stacked

Multi-Head Attention (1)

- **Single-head attention:** self-attention layer
- **Multi-head attention:** multiple *independent* self-attention layers
 - Output concatenated, then projected down to input size



Multi-Head Attention (2)

- Specifically, each attention head i has parameters $\mathbf{W}_i^Q$, $\mathbf{W}_i^K$ and $\mathbf{W}_i^V$
 - Since we have a projection layer at the end, dimensions of queries, keys and values can more freely change
 - We denote size of queries and keys by D_k (the same because dot product)
 - We denote size of values by D_v
 - Then $\mathbf{W}_i^Q \in R^{D \times D_k}$, $\mathbf{W}_i^K \in R^{D \times D_k}$ and $\mathbf{W}_i^V \in R^{D \times D_v}$ where D is size of input tokens
- As before, stack inputs of size D to form matrix $\mathbf{X}$
 - Then, $\mathbf{Q}_i = \mathbf{XW}_i^Q$, $\mathbf{K}_i = \mathbf{XW}_i^K$, $\mathbf{V}_i = \mathbf{XW}_i^V$
- We thus have **head** $_i = \text{Self-Attention}(\mathbf{Q}_i, \mathbf{K}_i, \mathbf{V}_i)$
- Outputs of each head is concatenated, projected down to D via $\mathbf{W}^O$
 - $\text{MultiHeadAttention}(\mathbf{X}) = (\text{head}_1 (+) \text{head}_2 (+) \dots (+) \text{head}_H) \mathbf{W}^O$, where $(+)$ denotes concatenation


Multi-Head Attention (3)

- Perhaps more important: **why?**
 - What is the **intuition behind this?**
 - How can we think about multi-head attention?
- **Different words in a sentence relate** to each other in **different ways**
 - *“They happily played Mario Kart until the sun came out.”*
 - *“Mario Kart”* is the object of verb played
 - *“Mario Kart”* described further by indirect complement *“until the...”*, etc.
- A single-head attention must capture all of these dependencies
 - A **multi-head attention splits the workload into independent heads**
 - Think of tasking a set of people to do one job, e.g. memorize a book
- Thus, multi-head attention is more powerful
 - But more difficult to train (more parameters)
- **Additional intuition:** multiple convolution filters in convolution layer
 - Similar principle: many patterns to capture, multiple filters allows model to capture different patterns using different filters

Positional Encoding (1)

- In the transition from RNNs to transformers, we lost something!
 - **We no longer “see” relative positions (order) of each token in the input sequence**
 - Each token attends to each other token, wherever they are
- Think about a self-attention layer
 - It produces contextualized representations for each input token
 - If we **shuffle the input sequence**, the output of the **linear combination** that produces contextualized representations is **still the same**
- But relative position of tokens matters in a sequence!
- **Solution:** modify input embeddings to encode position before applying attention
 - For example, learn position embeddings just as you do word embeddings
 - E.g. embedding for position 1, 2, etc.
 - Then combine embeddings of positions with words before multi-head attention, e.g. by adding them up -> $\mathbf{w}_1 + \mathbf{w}_{the}$

Positional Encoding (2)



- **Other types of positional encodings**, e.g. non-parameterized ones
 - We'll see some in detail in tutorials

Layer Normalization


- **Goal:** normalize output of any layer in a network
 - Not exclusive to transformers, commonly used in deep learning models
 - Improves gradient-based training
- It's a form of centering the data
 - Given output vector $\mathbf{x}$ for some hidden layer, compute its mean μ and standard deviation σ
 - Then center $\mathbf{x}' = (\mathbf{x} - \mu) / \sigma$
- Final value is actually $\text{LayerNorm} = \gamma \mathbf{x}' + \beta$
 - Both γ, β are learnable parameters
- Layer normalization can play **important role during training**
 - First proposed by [Ba et al. \(2016\)](#)
 - Subsequent studies have focused on its impacts
 - Recent architectures include more layer normalization in other parts
 - E.g. after input embeddings, before and after self-attention
 - More on this when discussing LLMs

Residual Connections

- Recall the **vanishing gradient** problem:
 - Most **operators have impact on the gradient that flows backward** during training
 - Generally, **input gradients are multiplied by** an operator-specific **Jacobian**
 - This factor **can make output gradient smaller than input gradient**
 - **With depth, gradients can thus vanish**
- Residual connections designed to address this issue ([He et al. 2015](#))
- **Main idea:** given operator, add its input back to its output
 - Addition operator has no impact on gradient, passes it backward unchanged! (Details in a future tutorial.)
 - Also not exclusive to transformers, commonly done in deep models
- In transformers, $\mathbf{z} = \text{LayerNorm}(\mathbf{X} + \text{Self-Attention}(\mathbf{X}))$
 - Thus, whatever impact *Self-Attention* has on the gradient, the full gradient is still preserved by the additive term $\mathbf{X}$ during backpropagation
 - We discuss this in a bit more detail in tutorials


Projection FNN

- Often described as **MLP**
- FNN used to transform output of multihead-attention
 - In reality, two linear projections, usually project up, then back down
 - ReLU (non-linear) activation between
- **Role of this component? Unclear**
 - Many different claims
 - Generally, typical Deep Learning component to give model freedom
 - Up projection -> more freedom
 - More later



Cross-Attention

- Important: **cross-attention**
 - Attention across different sequences
 - Note two attention layers in decoder
- 2nd multi-head attention in decoder (cross) attends to output of encoder
 - **Queries come from decoder tokens, keys/values from encoder tokens**
(see "encoder-decoder attention" in original transformer paper)
- Different transformer architectures
 - Encoder-decoder
 - Encoder-only, decoder-only
 - Cross-attention only present in some
 - More in lecture on LLMs



Dropout


- Proposed by [Srivastava et al. \(2014\)](#)
- **Idea:** randomly drop units in network (along with their connections)
 - It's a form of regularization, thus, only applied during training, not at inference time
 - Intuitively, network is dynamic during training
 - Prevents model from relying too much on static architecture by introducing variability, promoting generalization
- In practice, each units "zeroed" with probability p (hyperparameter)
 - E.g. see implementation in [PyTorch](#)
- **Original transformer used dropout** in two places
 - Specifically, before applying layer normalization and residual connection
 - That is, $\mathbf{z} = \text{LayerNorm}(\mathbf{X} + \text{Dropout}(\text{Self-Attention}(\mathbf{X})))$
 - Also after adding positional embeddings, i.e. $\mathbf{x} = \text{Dropout}(\mathbf{x}_1 + \mathbf{x})$
- Dropout still very common regularization method in LLMs

Inductive Bias in Transformers

- **Inductive bias:** set of assumptions made by a learning model
- Common story: self-attention is main innovation
 - But many more components in transformers
- **Transformer block:**
 - Two main components: **self-attention** and **MLP**
 - With corresponding layer normalizations and residual connections
 - Let's see the role of each
- **Self-attention:**
 - Linear operations between different input tokens
 - I.e. there is "**inter-token**" communication
 - Linearities useful for parallelization, but not expressive for learning
- **MLP:**
 - Transformations applied at each token independently
 - I.e. "**Intra-token**" phase
 - Non-linearities introduced here for expressivity


Visualizing The Transformer Block

- A.k.a. **transformer layer**
 - **In:** sequence of tokens
 - **Out:** sequence of contextualized tokens
- **Red arrows represent residual connections**
- Note **layer norms missing**, different models use them in different places today
- **All parallelizable!** Each contextualized vector in attention is independent, each matrix product in MLP is parallelizable



Ablation

- There is **no fundamental theory behind the design of deep learning models**
 - No way to accurately predict how a model will behave given its architecture
 - No equation that tells us exactly which components a model should have to perform a given task
- **Thus, ablation studies are essential!**
 - **Ablation:** remove one component of the model at a time to test its impact
 - E.g. remove normalization layer, run model again, how does performance change?
 - Or remove/change positional encodings. What impact does this have in performance?
- With ablation, we may gain some understanding of why a model works



Summary

- **Self-attention:** main innovation in transformer architecture
 - Each input token contextualized based on entire input sequence
- Architecture has **many components**
 - Self-attention
 - Positional encodings
 - Layer normalization
 - FNN
- Architectures are **normally deep** (stacked transformer blocks)
- Different flavors
 - Encoder-decoder, decoder-only, etc.
- **Large language models (LLMs)** based on transformers, with variations
 - More layer normalizations
 - Different positional embeddings
 - Etc.

References

- Speech and Language Processing, Jurafsky et al., 2024
 - Chapter 10
- References linked in corresponding slides