

Advanced Methods in Text Analytics

State Space Models



An Alternative to Transformers (1)

- [illegible]

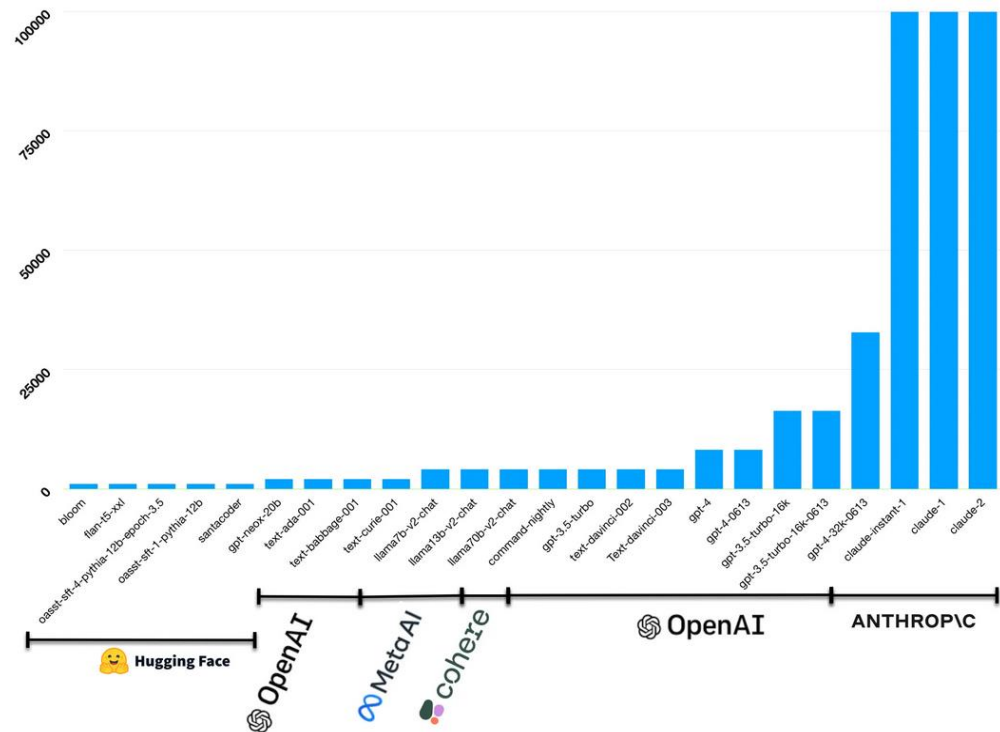


Image source

An Alternative to Transformers (2)

- Enter: state-space models
 - Approach from [control theory](#)
 - They are sequence models that do not use self-attention
 - Great at modeling very long-range sequences (e.g. 1M tokens)
- Today:
 - **State space models**, including S4 model (precursor to Mamba), introduced structure to state space models for efficiency, reached SOTA in long range dependency tasks
 - Relation to **RNNs** (you already know more about this than you think)
 - The **Mamba architecture**, recent state space model, showed competitive performance in causal language modeling when compared with LLMs twice its size (using less memory and compute)
- **Warning:**
 - Much of this comes from a **research area** that most NLP researchers (myself included) are **likely unfamiliar** with
 - **Some things** will be discussed in **detail**, **others** will be kept at **high-level**
 - As in the literature, many things illustrated with equations, not images

Outline

1. State Space Models
2. The Mamba Architecture

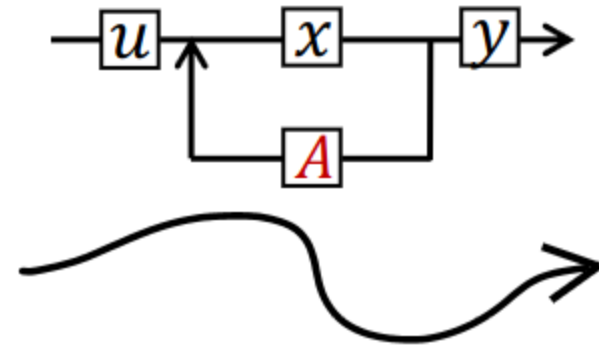
State Space Models

Continuous State Space Models (1)

- **Long range dependencies:** big challenge in sequence modelling
 - Materialized by **Long Range Arena** Benchmark ([Tay et al. 2021](#))
 - Wide range of tasks requiring processing between 1K and 16K tokens
- Most families of sequence models have variants to tackle this challenge
 - E.g. Lipschitz RNNs ([Erichson et al. 2021](#)), FlashAttention-2 ([Dao et al. 2023](#))
- Family of models that excel at this: **state space models (SSMs)**
- A ***continuous-time* latent state model** does the following:
 - 1. Maps 1-dimensional signal $x(t)$ to n-dimensional latent state $h(t)$
 - 2. Projects latent state (plus input signal) to 1-dimensional output signal $y(t)$
- Concretely:
 - 1. $dh(t)/dt = Ah(t) + Bx(t)$
 - 2. $y(t) = Ch(t) + Dx(t)$
- Here, **A , B , C , D** are learned parameters and t is time (usually $D = 0$ for simplicity because term $Dx(t)$ is seen as "skip connection" between input and output, easy to compute)

Continuous State Space Models (2)

- Let's note a few things about SSMs:
 - $\frac{dh(t)}{dt} = \mathbf{A}h(t) + \mathbf{B}x(t)$
 - $y(t) = \mathbf{C}h(t) + \mathbf{D}x(t)$
- The 1st equation is differential
 - Describes how \mathbf{h} changes over time
- The 2nd equation is **signal-to-signal model**
- Both $\frac{dh(t)}{dy}$ and $y(t)$ fully described by current input signal and hidden state
 - I.e. a **Markov decision process**
 - Ideally, state encodes all past history, enables us to determine future output $y(t)$
- Great, but this is NLP!
 - What are we supposed to do we these *continuous* SSMs?**
 - We are all about processing discrete sequences here, right?
- Yes, so **let's discretize SSMs**



$$\dot{x} = \mathbf{A}x + \mathbf{B}u$$

$$y = \mathbf{C}x + \mathbf{D}u$$

**Continuous
State Space**

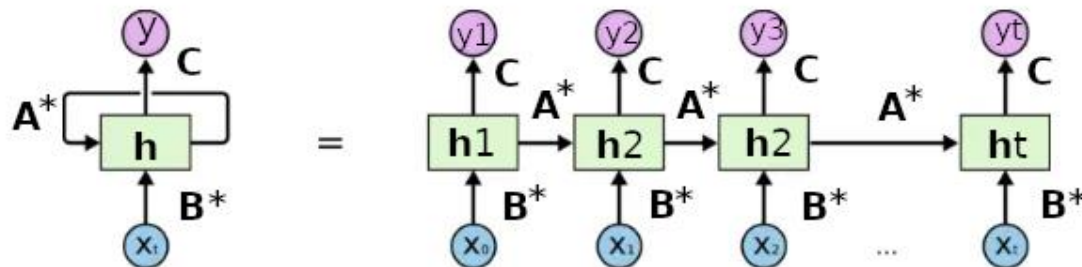
[Gu et al. 2022](#)

Discrete State Space Models (1)

- There are several ways to **discretize a continuous signal**
 - Think of it as **sampling a continuous signal at fixed time intervals**
 - Different SSMs use different approaches, each with different properties
- Simple way to get an *intuition for discretization*: [the Euler method](#)
- Let's start with our continuous model:
 - $\frac{dh(t)}{dt} = \mathbf{A}h(t) + \mathbf{B}x(t)$ (1)
 - $y(t) = \mathbf{C}h(t) + \mathbf{D}x(t)$ (2)
- 1st, we clean up notation a bit: $x_t = x(t)$, $\mathbf{h}_t = \mathbf{h}(t)$ and $\mathbf{h}_{t+1} = \mathbf{h}(t + \Delta)$
- 2nd, by definition of a derivative: $\frac{dh(t)}{dy} = (\mathbf{h}(t + \Delta) - \mathbf{h}(t)) / \Delta$ (3)
- Then, by plugging (3) in (1), we have:
 - $\mathbf{h}_{t+1} - \mathbf{h}_t = \Delta(\mathbf{A}\mathbf{h}_t + \mathbf{B}x_t)$
 - $\mathbf{h}_{t+1} = \Delta\mathbf{A}\mathbf{h}_t + \Delta\mathbf{B}x_t + \mathbf{h}_t$
 - $\mathbf{h}_{t+1} = \mathbf{h}_t(\Delta\mathbf{A} + \mathbf{I}) + \Delta\mathbf{B}x_t$
 - $\mathbf{h}_{t+1} = (\mathbf{I} + \Delta\mathbf{A})\mathbf{h}_t + (\Delta\mathbf{B})x_t$

Discrete State Space Models (2)

- If we now set $\mathbf{A}^* = (\mathbf{I} + \Delta\mathbf{A})$, $\mathbf{B}^* = \Delta\mathbf{B}$ and set $\mathbf{D} = \mathbf{0}$, we have:
 - $\mathbf{h}_{t+1} = \mathbf{A}^*\mathbf{h}_t + \mathbf{B}^*\mathbf{x}_t$ (1)
 - $\mathbf{y}_t = \mathbf{C}\mathbf{h}_t$ (2)
- Does that look familiar?
 - I hope we haven't forgotten! It looks like an RNN



An unrolled recurrent neural network.

[Image source](#)

- Discrete SSMs are essentially linear RNNs (i.e. seq2seq models)
 - I.e. RNN with activation $f = \text{I}$ when computing \mathbf{h} , i.e. $\mathbf{h}_{t+1} = f(\mathbf{A}^*\mathbf{h}_t + \mathbf{B}^*\mathbf{x}_t)$
 - Computing \mathbf{A}^* , \mathbf{B}^* (discretization) can be seen as first step in compute graph

Discrete State Space Models (3)

- As mentioned, different SSMs use different approaches for discretization
 - None actually uses the Euler method
- E.g. [S4](#) uses a bilinear method:
 - $\mathbf{A}^* = (\mathbf{I} - \Delta/2 \mathbf{A})^{-1}(\mathbf{I} + \Delta/2 \mathbf{A})$
 - $\mathbf{B}^* = (\mathbf{I} - \Delta/2 \mathbf{A})^{-1} \Delta \mathbf{B}$
- [Mamba](#) uses a process called [zero-order hold](#):
 - $\mathbf{A}^* = \exp(\Delta \mathbf{A})$
 - $\mathbf{B}^* = (\Delta \mathbf{A})^{-1}(\exp(\Delta \mathbf{A}) - \mathbf{I}) \Delta \mathbf{B}$
- Such models are thus parameterized by $\boldsymbol{\theta} = [\Delta, \mathbf{A}, \mathbf{B}, \mathbf{C}]$
- Ok, fine, so we are back to (linear) RNNs. Is that it?
- Not quite, as **RNNs fell out of favor for two specific reasons**
 - 1. Difficulty modeling long range dependencies
 - 2. Slow to train due to sequential computation of hidden states
- **SSMs need to deal with these issues** to be successful
 - Let's have a high-level look at how SSMs addressed these "RNN challenges"

Modeling Long Sequences with SSMs

- **HiPPO framework** introduced by [Gu et al. 2020](#)
 - **H**igh-order **P**olynomial **P**rojection **O**perators
 - Math-heavy work (described as "a bit of magic" by [some sources](#))
- Essentially, framework that proposes a **principled way of initializing matrix A** so state $h(t)$ can memorize input sequence x
 - Specifically, a complex matrix in the following upper triangular form:

$$(\text{HiPPO Matrix}) \quad A_{nk} = - \begin{cases} (2n+1)^{1/2}(2k+1)^{1/2} & \text{if } n > k \\ n+1 & \text{if } n = k \\ 0 & \text{if } n < k \end{cases}$$

- More generally, they formalized notion of *memory as online function approximation* (based on [approximation theory](#) and [signal processing](#))
 - The framework derived [GRUs](#) and [LMUs](#) from first principles
- Their results put SSMs on the map, reported **massive improvements**
 - E.g. from 68% to 98% on sequential MNIST benchmark
- Still, SSMs were expensive to train compared to Transformers
 - One RNN challenge down, one to go

Training SSMs Efficiently

- Important property of SSMs: **linear time invariance** (LTI)
 - Well known connection with convolutions (yet no citation in papers!)
- In short, because computing \mathbf{h}_{t+1} is linear, we can unroll SSMs as follows:
 - $\mathbf{h}_0 = \mathbf{B}^* \mathbf{x}_0$ $\mathbf{h}_1 = \mathbf{A}^* \mathbf{B}^* \mathbf{x}_0 + \mathbf{B}^* \mathbf{x}_1$ $\mathbf{h}_2 = \mathbf{A}^{*2} \mathbf{B}^* \mathbf{x}_0 + \mathbf{A}^* \mathbf{B}^* \mathbf{x}_1 + \mathbf{B}^* \mathbf{x}_2$
 - $\mathbf{y}_0 = \mathbf{C} \mathbf{B}^* \mathbf{x}_0$ $\mathbf{y}_1 = \mathbf{C} \mathbf{A}^* \mathbf{B}^* \mathbf{x}_0 + \mathbf{C} \mathbf{B}^* \mathbf{x}_1$ $\mathbf{y}_2 = \mathbf{C} \mathbf{A}^{*2} \mathbf{B}^* \mathbf{x}_0 + \mathbf{C} \mathbf{A}^* \mathbf{B}^* \mathbf{x}_1 + \mathbf{C} \mathbf{B}^* \mathbf{x}_2$
 - and so on, where initial state $\mathbf{h}_{-1} = \mathbf{0}$.
- Thus, for input sequence \mathbf{x} , the entire output sequence is $\mathbf{y} = \mathbf{K} \odot \mathbf{x}$ where \odot is the (discrete) **convolution** operation and \mathbf{K} is the following (giant) convolution filter:
 - $\mathbf{K} = (\mathbf{C} \mathbf{B}^*, \mathbf{C} \mathbf{A}^* \mathbf{B}^*, \mathbf{C} \mathbf{A}^{*2} \mathbf{B}^*, \dots, \mathbf{C} \mathbf{A}^{*|\mathbf{x}|-1} \mathbf{B}^*)$
- I.e., **we can compute entire forward pass as single global convolution!**
 - Can be computed efficiently with [Fast-Fourier Transforms](#) (FFTs) once \mathbf{K} is known
 - But computing \mathbf{K} is expensive (repeated matrix multiplication by \mathbf{A}^*)
- Achieving this in linear time was a main contribution of the S4 model

Structured State Space Models

- **S4**: structured state space sequence models ([Gu et al. 2022](#))
 - Again, math-heavy paper, only high-level discussion here
- **Main contribution**:
 - Computing global kernel K in linear time (normally quadratic)
- They **enforce structure** to parameter matrix A in SSMs
 - A is special form of **diagonal matrix**: diagonalizable by normal matrices
 - Work far from trivial, as matrix A had to retain the properties of the HiPPO matrix that accounts for the success of SSMs at long range dependencies
- An S4 layer was proposed to construct deep S4 models
 - **S4 block**: SSM seq2seq model + dropout + non-linearity + linear projection
 - Since S4 processes single scalar, **S4 block contains H copies of S4 model**
 - Then, each block processes inputs of size (batch size, input length, hidden size), same as RNNs, Transformers (but S4 is linear in input length!)
 - Can be **computed as convolution** (training) or **autoregressively** (inference) (nicely shown in code [here](#))
- So, with S4 we get **efficient SSM training without performance loss**
 - Now let's recap SSMs vs RNNs before we look at Mamba

SSMs vs RNNs

- While similar, there are **important differences between SSMs and RNNs**
 - 1. SSMs (and linear RNNs) can be efficiently parallelized (global convolution)
 - 2. SSMs are "linear RNNs" but with special requirements for how to compute some parameters (discretization)
 - 3. SSMs are complex valued and initialized according to the HiPPO theory
- Also, note that while SSM matrices \mathbf{A}^* , \mathbf{B}^* are equivalent to RNN parameters, they actually share parameters themselves (usually Δ and \mathbf{A})
- In short, **these differences are important**, account for success of SSMs!
- After S4's success, **recent work "resurrected RNNs"** ([Orvieto et al. 2023](#))
 - Started with a linear RNN, "ablated" their way to reach S4's performance
 - Specifically, they introduced **diagonalization**, **special initialization** and **parameterization**, as well as **normalization** to RNNs
 - Insightful in terms of where success of SSMs comes from
- Having "covered" how both RNN challenges were addressed by SSMs, **we can now discuss Mamba!**
 - This because the model is an extension of the S4 architecture

The Mamba Architecture

Motivation

- Perspective proposed by authors:
 - *"A fundamental problem of sequence models is compressing context into a small state"*
- For example, **self-attention** is:
 - **Effective:** does not compress context at all, stores it entirely (KV cache in causal LMs)
 - **Inefficient:** KV cache takes linear space in input sequence length
- Conversely, **RNNs** are:
 - **Efficient:** compress context in hidden state h (size usually much smaller than the long range sequences we want to model with SSMs)
 - **Ineffective:** especially for long sequences, as performance depends on how well they compress context (though the "RNN resurrection" paper improved considerably on this thanks to research in SSMs)
- In short, **efficient vs effective trade-off** summarized by how well a model compresses their state
 - Proposed **fundamental principle** for building sequence models: **selectivity**

Improving SSMs with Selection

- **Selectivity:** context-aware ability to focus on or filter out inputs into a sequential state
 - Sounds awfully familiar, right?
- In other words, an **attention-like mechanism** that **controls how information propagates** or **interacts** between elements of a sequence
 - Is it really an attention-like mechanism?
- Yes! The **main Mamba design innovation** is allowing **more access to input sequence** by letting the **SSM parameters** be **input-dependent**
 - Recall attention on RNNs was basically: *Hey, why don't we just access the input at each inference step instead of just interacting with the hidden state?*
 - Mamba is the SSM that asks the same question (likely inspired by the success of the Transformer).
- Specifically, some parameters became functions of input sequence \mathbf{x}
 - Matrix $\mathbf{B} = \mathbf{W}_B \mathbf{x}$, matrix $\mathbf{C} = \mathbf{W}_C \mathbf{x}$ where \mathbf{W}_B , \mathbf{W}_C are learned projections
- **Problem:** this change means we lose the ability to train SSMs efficiently
 - Let's see why

Improving SSMs with Selection?

- Recall efficient training of SSMs:
 - Given input sequence \mathbf{x} , entire output sequence is $\mathbf{y} = \mathbf{K} \odot \mathbf{x}$ where \odot is (discrete) convolution operation and \mathbf{K} is the following (giant) convolution kernel: $\mathbf{K} = (\mathbf{CB}^*, \mathbf{CA}^*\mathbf{B}^*, \mathbf{CA}^{*2}\mathbf{B}^*, \dots, \mathbf{CA}^{*|\mathbf{x}|-1}\mathbf{B}^*)$
 - Efficient! Compute kernel elements separately, apply [convolution theorem](#)
- Now, note the following:
 - \mathbf{h}_t depends on \mathbf{h}_{t+1} because $\mathbf{h}_{t+1} = \mathbf{A}^*\mathbf{h}_t + \mathbf{B}^*\mathbf{x}_t$
 - Since $\mathbf{B} = f(\mathbf{x})$, $\mathbf{B}^* = f(\mathbf{x})$ because \mathbf{B}^* is a function (discretized form) of \mathbf{B}
 - For $t = 1$, $\mathbf{x} = [\mathbf{x}_1]$, for $t = 2$, $\mathbf{x} = [\mathbf{x}_1, \mathbf{x}_2]$, for $t = n$, $\mathbf{x} = [\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_n]$
- So, \mathbf{B}^* is now time dependent, i.e. \mathbf{B}^*_t
 - Thus, a more accurate notation now is: $\mathbf{h}_{t+1} = \mathbf{A}^*\mathbf{h}_t + \mathbf{B}^*_t\mathbf{x}_t$ since \mathbf{B}^* is a function of \mathbf{x} , which changes per time step (autoregressive models)
- In other words, **hidden states now need to be computed in sequence**
 - We lost one of the two improvements SSMs had over RNNs
 - So, another main **contribution from Mamba: make training efficient again**

Hardware-Aware Selectivity

- To regain training efficiency, the authors did two things:
 - 1. Switch from convolution to a [scan](#) operation
 - 2. Implement **scan in a GPU-optimized way**
- **Scan:** otherwise known as a cumulative sum
 - E.g. given sequence $x = [3, 2, 6]$, output is sequence $y_1 = 3, y_2 = 5, y_3 = 11$
 - Can be parallelized to be computed in $\log n$ (nicely shown [here](#))
 - Already used for efficiently training linear RNNs ([Orvieto et al. 2023](#))
- So, in short: **what is Mamba?**
 - An **S4 model** (structured SSM) with **Selectivity**
 - An S4 model trained with the **Scan** operation
 - **Hence**, authors described it as an **S6 model**
- **Transformers vs Mamba:** [overall training costs](#) for sequence of length n
 - **Memory:** $O(n)$ vs $O(1)$ (KV cache vs hidden state)
 - **Runtime:** $O(n^2)$ vs $O(n)$ (self-attention vs GPU-optimized SSM)
- In other words, massive improvements!
 - If performance on par with LLMs, we may have a new king

The Mamba Block

- As with the Transformer, we have a **Mamba block**

- **Inter-token communication:**

- Transformer: self-attention
- Mamba: SSM

- **Intra-token computation:**

- Transformer: MLP
- Mamba: MLP

- **Efficiency:**

- Transformer: linear attention
- Mamba: linear RNN

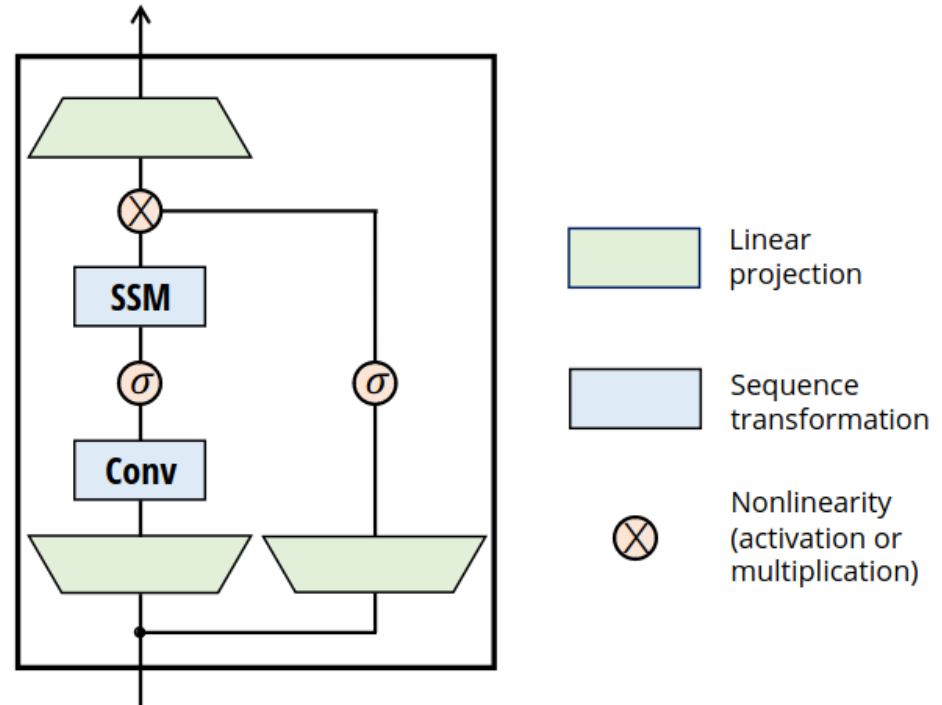
- **Expressivity:**

- Transformer: non-linearity between MLP projections
- Mamba: non-linearity between MLP projections

- In practice:

- Up-projection by a factor of 2 (compared to usual 4 in LLMs)

- [Silu/Swish](#) activations, optional layer normalization before/after block



[Gu et al. 2023](#)

Mamba vs LLMs

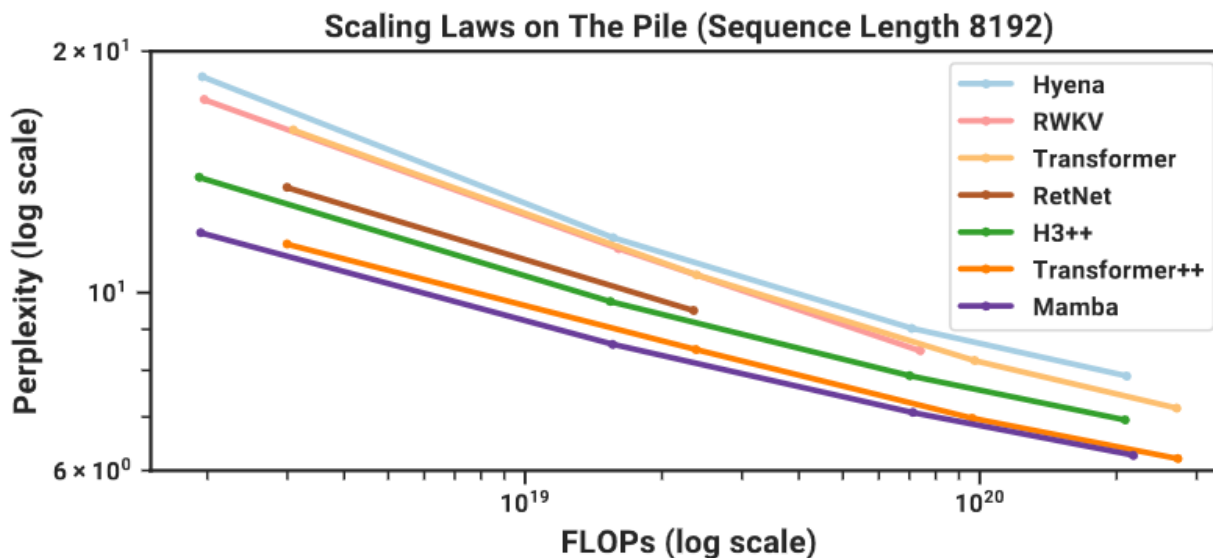
- Some **results** on **downstream tasks** (with **5x faster inference** compared to transformer-based LLMs)

MODEL	TOKEN.	PILE PPL ↓	LAMBADA PPL ↓	LAMBADA ACC ↑	HELLASWAC ACC ↑
GPT-Neo 2.7B	GPT2	—	5.63	62.2	55.8
Hybrid H3-2.7B	GPT2	—	7.92	55.7	59.7
OPT-2.7B	OPT	—	5.12	63.6	60.6
Pythia-2.8B	NeoX	6.73	5.04	64.7	59.3
RWKV-3B	NeoX	7.00	5.24	63.9	59.6
Mamba-2.8B	NeoX	6.22	4.23	69.2	66.1
GPT-J-6B	GPT2	—	4.10	68.3	66.3
OPT-6.7B	OPT	—	4.25	67.7	67.2
Pythia-6.9B	NeoX	6.51	4.45	67.1	64.0
RWKV-7.4B	NeoX	6.31	4.38	67.2	65.5

- Some **concerns** from the community: [Gu et al. 2023](#)
 - Are baselines strong enough? E.g. no latest OpenAI models
 - More importantly, **why no comparison with larger LLMs?**

Scaling Laws

- Authors also provided scaling laws on all tested tasks (here only LM)
 - They followed the [Chinchilla protocol](#)
 - Roughly, statistical model that measures loss as function of model size



[Gu et al. 2023](#)

- **"First attention-free model to match performance of very strong Transformer recipe (Transformer++)"** (with cheaper memory/compute!)
 - Transformer++ recipe: Rotary, SwiGLU, MLP, RMSProp, no linear bias, higher learning rate (PaLM, LLaMA)

A Different Paradigm (1)

- Nice discussion on a potential paradigm shift by [Kola Ayonrinde](#)
 - Specifically about using pre-trained causal language models (CLMs)
- **Sources of information** in transformer-based CLMs **at inference time?**
 - **Training data**, i.e. pre-trained weights (long-term memory, but compressed)
 - **In-context data**, i.e. prompt (short-term memory, must be read every time)
- Do we have **selectivity when using such CLMs?**
 - Yes, via **prompting** (what to leave in or out), **RAG**, etc.
 - This because entire prompt is accessed (attention) to predict next tokens
- What about when using **SSM-based CLMs like Mamba?**
 - Training data is there in the same way, i.e. compressed long-term memory
 - Selectivity is a core design of the system as it reads any input
 - But **in-context data is also long-term now!** Why?
- Because **we can store the last hidden state, plug it in later as needed!**
 - Transformers not designed to have a state or representation plugged in
 - They must instead "read" the entire input to produce contextualized vectors

A Different Paradigm (2)

- This is **essentially a new form of prompting**
 - Swapping hidden states as needed
 - A modularity akin to swapping LoRA modules
- Remember: **SSMs are very good at long range dependencies**
 - We could potentially read entire set of textbooks, store hidden state h_n
 - Then, at inference time, use h_n as initial state of our prompt
- **Question:** can transformers' contextual embeddings also act as states?
 - There is already [some work](#) on this, so maybe...
- And with that, our discussion on SMMs comes to an end
 - So, let's summarize!

Summary

- **State Space Models (SSMs):**
 - **Sequence models similar to linear RNNs, but with key differences**
 - Differences make them **very good at long range dependencies**
 - This ability comes from **principled design choices** (HiPPO framework)
- **Mamba Architecture:**
 - **SSM with attention-like mechanism** for selectively using input sequence
 - Performs **competitively with LLMs twice its size**
 - Scaling laws show promise for scaling up the model
- **SSMs vs Transformers**
 - **SSMs much more efficient**, memory: $O(1)$ vs $O(n)$, compute: $O(n)$ vs $O(n^2)$
 - But **no results on Mamba vs largest LLMs yet**, e.g. 70B or larger

References

- [Efficiently Modeling Long Sequences with Structured State Spaces](#) by Gu et al, 2022
- [The Annotated S4](#) by Sasha Rush and Sidd Karamcheti
- [Resurrecting Recurrent Neural Networks for Long Sequences](#) by Orvieto et al. 2023
- [Mamba: Linear-Time Sequence Modeling with Selective State Spaces](#) by Gu and Dao, 2023
- [Mamba Explained](#) by Kola Ayonrinde
- [Mamba from Scratch](#) by Algorithmic Simplicity
- References linked in corresponding slides