

Advanced Methods in Text Analytics

Exercise 6: Transformers - Part 2

Solutions

Daniel Ruffinelli

Licensed under Creative Commons Attribution 4.0 International

FSS 2026

1 The Residual Stream

(a) We have:

$$\text{SA}(\mathbf{X}) = \text{softmax} \left(\frac{\mathbf{X}\mathbf{W}^Q(\mathbf{X}\mathbf{W}^K)^T}{\sqrt{d}} \right) \mathbf{X}\mathbf{W}^V \quad (1)$$

where the softmax function is applied row-wise and $\mathbf{X}\mathbf{W}^Q = \mathbf{Q}$ is referred to as the query matrix, $\mathbf{X}\mathbf{W}^K = \mathbf{K}$ is the key matrix, and $\mathbf{X}\mathbf{W}^V = \mathbf{V}$ the value matrix. Recall that the factor $\frac{1}{\sqrt{d}}$ is used for numerical stability (scaled dot-product attention).

(b) We have:

$$\text{FNN}(\mathbf{X}) = \text{ReLU}((\mathbf{X}\mathbf{W}_1 + \mathbf{b}_1)) \mathbf{W}_2 + \mathbf{b}_2 \quad (2)$$

where parameters are $\mathbf{W}_1 \in \mathbb{R}^{d \times m}$, $\mathbf{b}_1 \in \mathbb{R}^m$, $\mathbf{W}_2 \in \mathbb{R}^{m \times d}$, and $\mathbf{b}_2 \in \mathbb{R}^d$.

(c) We do this in stages. First, we focus on SA, its corresponding residual connection and layer normalization LN_1 . We have:

$$\mathbf{O}_1 = \text{LN}_1(\text{SA}(\mathbf{X}) + \mathbf{X}) \quad (3)$$

Then, we focus on FNN, its corresponding residual connection and layer normalization area applied. That is:

$$\mathbf{O}_2 = \text{LN}_2(\text{FNN}(\mathbf{O}_1) + \mathbf{O}_1) \quad (4)$$

All together, we have:

$$\text{TF}(\mathbf{X}) = \text{LN}_2(\text{FNN}(\text{LN}_1(\text{SA}(\mathbf{X}) + \mathbf{X})) + \text{LN}_1(\text{SA}(\mathbf{X}) + \mathbf{X})). \quad (5)$$

(d) Layer normalization is used to normalize the vectors it is given as input. Concretely, given a vector $\mathbf{x} \in \mathbb{R}^d$, the layer normalization operation is given by:

$$\text{LN}(\mathbf{x}) = \frac{\mathbf{x} - \mu}{\sigma} \gamma + \beta \quad (6)$$

where γ and β are learned parameters. This operation can be seen as a form of data centering and is commonly used so keep training stable, sometimes at the cost of performance

on some downstream tasks. It can be applied anywhere in a network, including anywhere in a transformer layer. In fact, a common variant is the *pre-norm* transformer, which applies layer normalization *before* each of SA and FNN. Typically, this variant also includes a single final layer normalization operation applied only to the output of the last transformer layer.

- (e) To update parameters θ_{SA} of SA w.r.t. some loss function L during training, we need:

$$\frac{\partial L}{\partial \theta_{SA}} = \frac{\partial L}{\partial SA(\mathbf{H})} \frac{\partial SA(\mathbf{H})}{\partial \theta_{SA}}, \quad (7)$$

where $\mathbf{H}$ is the input to SA, which in turn is the output of the previous transformer layer TF_{i-1} . The term $\frac{\partial L}{\partial SA(\mathbf{H})}$ is the partial gradient received during backpropagation from the upper layer, the FNN operator in the context of a transformer layer.

- (f) Just as we received a gradient from the upper layer during training, we pass down the following gradient to the lower transformer layer TF_{i-1} during backpropagation:

$$\frac{\partial L}{\partial \mathbf{H}} = \frac{\partial L}{\partial SA(\mathbf{H})} \frac{\partial SA(\mathbf{H})}{\partial \mathbf{H}}. \quad (8)$$

This is the gradient of the output of TF_{i-1} w.r.t. L , and is needed to compute the weight updates of the operators in lower layers that contribute to the computation of the input to TF_i , i.e. $\mathbf{H}$.

- (g) Let $\mathbf{O}_1$ be the output of SA with a residual connection around it, i.e. $\mathbf{O}_1 = SA(\mathbf{H}) + \mathbf{H}$. Then, the gradient “received” from the upper layer is $\frac{\partial L}{\partial \mathbf{O}_1}$. Eq. 8 can then be rewritten as follows:


$$\begin{aligned} \frac{\partial L}{\partial \mathbf{H}} &= \frac{\partial L}{\partial \mathbf{O}_1} \frac{\partial \mathbf{O}_1}{\partial \mathbf{H}} \\ &= \frac{\partial L}{\partial \mathbf{O}_1} \frac{\partial (SA(\mathbf{H}) + \mathbf{H})}{\partial \mathbf{H}} \\ &= \frac{\partial L}{\partial \mathbf{O}_1} \left(\frac{\partial SA(\mathbf{H})}{\partial \mathbf{H}} + \frac{\partial \mathbf{H}}{\partial \mathbf{H}} \right) \\ &= \frac{\partial L}{\partial \mathbf{O}_1} \frac{\partial SA(\mathbf{H})}{\partial \mathbf{H}} + \frac{\partial L}{\partial \mathbf{O}_1}. \end{aligned}$$

In other words, when using residual connections, the gradient $\frac{\partial L}{\partial \mathbf{O}_1}$ we receive from higher layers during backpropagation is passed down to lower layers by making a modification via the addition operation.

Note that without the residual connection, the gradient we pass down would be modified by the multiplying factor $\frac{\partial SA(\mathbf{H})}{\partial \mathbf{H}}$, and the additive term would not exist. Since such a multiplying factor can quickly reduce the size of the gradient that is passed backward during training, the residual connection was designed by [He et al. 2015](#) so that gradients can be more safely passed through an operator without it having a severe impact on those gradients.

The name “residual” connection comes from the fact that, due to this change to prevent gradients from vanishing during training, the function $f(\mathbf{X}) = \text{OP}(\mathbf{X})$ computed by some operator becomes $f(\mathbf{X}) = \text{OP}(\mathbf{X}) + \mathbf{X}$. This means that we are learning to apply a suitable *difference* $\text{OP}(\mathbf{X})$ that is added to given input $\mathbf{X}$. In the context of a transformer

layer, $\mathbf{X}$ is a set of contextualized representations, so what each transformer layer learns is to add something to it *if needed*. If $\mathbf{X}$ is already suitable for the task, then a layer may make small modifications to it. This suggests that the main flow of information is coming from $\mathbf{X}$, and not from $\text{OP}(\mathbf{X})$, or in other words, from the residual connections and not from the transformer layers. This perspective is known as the “residual stream” view of transformers, illustrated in the image below from Jurasfky and Martin (2024), Section 10.4. For evidence of the internal mechanisms of transformers being well described by this perspective, see [here](#) and [here](#).



The Residual Stream. Information flows through residual connections.

2 Attention Heads are Independent and Additive

(a) We have:

$$\text{MHA}(\mathbf{X}) = (\text{SA}_1(\mathbf{X}) \oplus \text{SA}_2(\mathbf{X}) \oplus \dots \oplus \text{SA}_k(\mathbf{X})) \mathbf{W}^O + \mathbf{X}, \quad (9)$$

where $\oplus$ is the concatenation operator, $\mathbf{W}_O \in \mathbb{R}^{kd \times d}$ is a parameter matrix that projects the concatenated vector back to d -dimensional residual space, and the additive term is the residual connection around MHA. Note that the size of $\mathbf{W}^O$ is, indirectly, a design choice, as it depends on the number of attention heads (in this case k) and the size of tokens (in this case d).

(b) The operation in question here is the projection to residual space done by $\mathbf{W}^O$. To see how this is parallelizable w.r.t. each $\text{SA}_i(\mathbf{X})$, let $\mathbf{O}^i \in \mathbb{R}^{n \times d}$ be the output of SA_i , as described. After concatenating the output of each attention head, we have $\mathbf{O} \in \mathbb{R}^{n \times kd}$ where k is the number of attention heads. That is, we have the following block matrix:

$$\mathbf{O} = [\mathbf{O}^1 \mathbf{O}^2 \dots \mathbf{O}^k]. \quad (10)$$

Now, note that when computing each element ij of $\mathbf{O}\mathbf{W}^O$, i.e. $[\mathbf{O}\mathbf{W}^O]_{ij}$, we compute the following dot product:

$$[\mathbf{O}\mathbf{W}^O]_{ij} = \sum_{h=1}^{kd} \mathbf{O}_{ih} \mathbf{W}_{hj}^O. \quad (11)$$

We can rewrite this dot product to make use of the $\mathbf{O}^i$ components of $\mathbf{O}$ more explicitly, as follows:

$$[\mathbf{O}\mathbf{W}^O]_{ij} = \sum_{h=1}^d \mathbf{O}_{ih} \mathbf{W}_{hj}^O + \sum_{h=d+1}^{2d} \mathbf{O}_{ih} \mathbf{W}_{hj}^O + \dots + \sum_{h=(k-1)d+1}^{kd} \mathbf{O}_{ih} \mathbf{W}_{hj}^O \quad (12)$$

$$= \sum_{h=1}^d \mathbf{O}_{ih}^1 \mathbf{W}_{hj}^{O_1} + \sum_{h=1}^d \mathbf{O}_{ih}^2 \mathbf{W}_{hj}^{O_2} + \dots + \sum_{h=1}^d \mathbf{O}_{ih}^k \mathbf{W}_{hj}^{O_k}, \quad (13)$$

where we now defined $\mathbf{W}^{O_i} \in \mathbb{R}^{d \times d}$ as the i -th block of $\mathbf{W}^O$ that corresponds to $\mathbf{O}^i$. That is, we can also see $\mathbf{W}^O$ as the following block matrix:

$$\mathbf{W}^O = [\mathbf{W}^{O_1} \mathbf{W}^{O_2} \dots \mathbf{W}^{O_k}]^\top. \quad (14)$$

- (c) From the solution to the previous subtasks, it follows that we can represent the general operation of MHA as:

$$\text{MHA}(\mathbf{X}) = [\mathbf{O}^1 \mathbf{O}^2 \dots \mathbf{O}^k] \begin{bmatrix} \mathbf{W}^{O_1} \\ \mathbf{W}^{O_2} \\ \dots \\ \mathbf{W}^{O_k} \end{bmatrix} + \mathbf{X}. \quad (15)$$

And being as this is a matrix multiplication of block matrices, we can write this as the following linear combination:

$$\text{MHA}(\mathbf{X}) = \sum_{i=1}^k \mathbf{O}^i \mathbf{W}^{O_i} + \mathbf{X}, \quad (16)$$

where $\mathbf{M}^i = \mathbf{W}^{O_i}$ as desired.

This equivalence shows that the output of the multi-head attention layer MHA can be expressed as a linear combination of the outputs of each attention head $\text{SA}_i(\mathbf{X})$, where the weights of that linear combination are learned by the layer. Thus, each (weighted) attention head contributes to the learned representations by *independently adding* to the representations in residual space. Indeed, the field of interpretability has identified that certain attention heads learn specific roles that completely differ from other attention heads. For more details, see [here](#).

3 Language Models with Transformers

- (a) Storing models during or after training is important to later be able to use them, e.g. for inference, fine-tuning or to continue training. PyTorch models have a state dictionary (state_dict) that is usually stored on disk in so-called *checkpoint* files. It's important that these files contain all the necessary information so that we may later be able to use the model. *register_buffer* is used to define “parameters” of a model that need to be stored in its state_dict, but that aren't actually parameters in the learning sense. That is, they are not changed during training. Since we are implementing a non-parameterized positional encoding, we register these embeddings as buffers, because we later need them for using the model.

- (b) The answer is not straightforward. If we look at the class from PyTorch that we use, it's the `TransformerEncoder` class. But, as we'll see later, we apply a mask to our attention scores, which have an impact on the type of model we are using, so we'll come back to this in the next few questions. In addition, the model does not have cross-attention (attention between decoder and the outputs of the encoder), so it is definitely not an encoder-decoder model either.

The code does have a component called a decoder, but what is that exactly? It's a linear layer that projects down to a vector space the size of our vocabulary, so while common, this naming convention can be confusing, as this decoding step refers to decoding in the tokenization sense. This projection, in combination with a softmax, will be used to produce probabilities over our vocabulary in order to predict the next token (if we construct the training examples correctly). But why isn't the softmax function there then? Well, as in our previous coding exercise, we will use `CrossEntropyLoss`, which includes the softmax function, so we need our model to output logits.

- (c) This function is used to construct an attention mask. Masking is the process by which we "block" some entries in a tensor. In this context, the forward pass of the `TransformerEncoder` takes a mask as input along side the input sequence. This mask is added to the matrix of attention *scores* in the self-attention layer. If scores are *-Inf*, their attention weight (via softmax) is zero. Entry *ij* in this matrix can be seen as how much attention input token *i* pays to *j*. Thus, the mask is a matrix with zeros in the lower triangular, and the rest *-Inf*.
- (d) Given the attention mask we use, as explained in the answer above, this is a causal language model. Note that the use of this *causal mask*, in combination with the lack of cross-attention, means we have a decoder-only model, despite using a `TransformerEncoder` class. It is common to implement such a CLM using `TransformerEncoder`.
- (e) See Jupyter notebook.
- (f) We get validation perplexity (PPL) of about 500, which is somewhat higher than with the RNN, which was about 400. Recall that PPL can be interpreted as proportional to the vocabulary size, which is about 30K. In addition, and as discussed in previous tutorials, it can be interpreted as the average *branching factor*, i.e. the average number of possible next words after a given one.

The PPL values from the RNN and this transformer are indeed comparable, as both models are evaluated on the same task using the same validation data (it's virtually the same code base). As for resources, they are also comparable. The embedding size is the same as the one used by the RNN, and both use dropout with default values as the only form of regularization. It's therefore likely that this slightly lower performance is due to the stronger inductive bias in the transformer model, i.e. it's overfitting. It would be nice to try to use a transformer with lower number of attention heads, or lower number of layers. And indeed, such tests are important to do on validation data before deciding which model will be used on test data or real-world data. We look at this further in the next question.

- (g) See Jupyter notebook for details. When training longer, for 20 epochs, we see the transformer achieving a PPL of about 400, i.e. it achieves the same performance achieved by the RNN in a single epoch. But when training the RNN for 20 epochs, it gets a PPL of about 350, so while smaller, there is still a performance gap in favor of the RNN.

W.r.t. changing model depth, we get slightly lower performance with a single layer, and a marginal improvement with 3 layers. But when using 6 layers performance drops, suggesting overfitting.

Finally, we could use a [learning rate scheduler](#) to get an increase in performance, as the learning rate is often a crucial hyperparameter in deep learning models, so processes like a [learning rate warmup](#) are common practice when training large language models.

In general, the transformer does not quite achieve the performance of the RNN in our tests, but it's possible that with other modifications we can get our transformer to performn better, e.g. using smaller hidden sizes, a different number of attention heads, etc. In addition, note that model stability should ideally be checked. In other words, how large is the variance of PPL when training the same model multiple times? Running the models in these tests multiple times suggests they are not quite stable, so that it's indeed possible that the transformer and the RNN are indeed performing similarly when comparing the mean and variance of PPL over, say, 10 different runs of each model.